

Glossary of Tech & Popular Terms

A shared vocabulary for community moderation in Wales

A community-facing reference that explains the words a moderation steward runs into — the technology behind the tool, the language of platforms and moderation, the tactics of online harm, the coded ways hostility hides, and the everyday internet slang that surrounds it all. It exists so that no one has to pretend to understand a term to take part in governing the system. It supports the **router’s glossary look-up task** and the prompt-level requirement to **use context-appropriate terminology and provide a glossary for unfamiliar terms**, and it serves the **community stewards** at Race Council Cymru and the Chinese in Wales Association directly.

Document	Glossary of Tech & Popular Terms — Corpus Material 3 of 4
Maps to	Layer 2 router (glossary look-up task) & Layer 3 prompt constraint (terminology / provide a glossary)
Serves	The LPA decision-support agent and community stewards (AI-literacy & shared vocabulary)
Companions	(1) Platform Terms & Policies; (2) Political & Cultural Context
Status	Living reference — online language changes fast; stewards should add and retire terms

How to read the coded-hostility sections

Sections 04–05 explain how hostility disguises itself, including specific coded terms and symbols. They are written strictly for **recognition**, drawing on established anti-hate education. Two rules govern their use: coded language is **dual-use and context-dependent** — the same word can be innocent, reclaimed, or hostile depending on who uses it and how — so a match is a prompt to look closer, never proof of intent; and the **live, evolving specifics** belong in the community-curated evasion register and in specialist monitors, not in a static glossary. Nothing here should be used to accuse an individual on a single signal.

Copyright & safety note. All definitions are written in the compilers’ own words. The glossary does not decode terminology used to trade or access child sexual abuse material; such requests are refused, not defined.

00 How this document fits the corpus

The requirements gathered from the community workshop ask for a 'glossary of tech and popular terms' in two places: as a **look-up task the router should recognise and serve**, and as a **prompt-level behaviour** — the agent should use terminology the community is comfortable with and offer a plain-language definition whenever it uses, or is asked about, an unfamiliar term. This document is the reference behind both.

It does three jobs at once:

- **AI literacy for stewards.** The community should be able to govern a system it understands. Section 01 translates the technology — including the plain-language metaphors used in the workshop — into terms a non-technical steward can hold onto.
- **Shared vocabulary for the task.** Sections 02–03 give common definitions of platform, moderation and online-harm terms so the agent and stewards mean the same thing by 'shadowban', 'brigading' or 'dog-whistle'.
- **Recognition of coded hostility.** Sections 04–05 explain the mechanics of evasion and give recognition aids, feeding the same concern as the community-curated evasion register (which holds the live specifics).

How entries are written. Each term is paired with a short, plain-language definition. Where a term is contested, dual-use, or sensitive, the definition says so rather than pretending it is neutral. Terms move in and out of use quickly online, so this is explicitly a living document for stewards to maintain.

01 The technology behind the tool

Plain-language definitions of the AI and platform technology a steward needs in order to understand, question and govern the moderation tool. The workshop introduced several of these through metaphors; those are noted so the two vocabularies line up.

Large language model (LLM)	The core AI behind the tool: a system trained on huge amounts of text to predict and generate fluent language. It is confident and articulate but has no built-in sense of truth — the workshop's 'confident stranger who never admits they don't know'.
Generative AI	AI that produces new content (text, images, audio) rather than only classifying existing content. A chatbot is one interface to it.
Hallucination	When an AI states something false or invented as if it were fact. The reason the framework insists the agent cite corpus entries and refuse when it has nothing to cite.
Prompt	The instructions or question given to an AI. The 'system prompt' is the standing instruction set that governs the agent's behaviour — the workshop's 'job description'.
Retrieval-Augmented Generation (RAG)	An approach where the AI must first look up approved reference documents (the 'retrieval corpus') and base its answer on them — the workshop's 'community filing cabinet the tool must check first'. This document is one item in that cabinet.
Retrieval corpus	The curated set of community-approved documents the tool draws on, including these corpus materials, the harm typologies and the evasion register. Inspectable and editable by stewards.
Router / classifier	Software that sorts each item into a category (for example pre-publication vs incoming, or by harm type) so the right rules apply — the workshop's 'specialist colleagues for different jobs'.
Toxicity score	A number an automated classifier assigns to estimate how 'toxic' a message is. The framework's analysis showed these scores routinely miss implicit and coded hostility, which is why they are not trusted on their own.
Fine-tuning	Retraining a model on extra data to change its behaviour. The framework deliberately avoids this in favour of editable corpus and rules, because fine-tuned models are opaque and hard for a community to inspect.
Training data	The text a model learned from. If it over-represents majority norms, the model inherits those biases — a root cause of the over- and under-moderation this project addresses.
Algorithm	In everyday use, any automated system that ranks, filters or moderates content. On platforms it decides what is shown, hidden or removed; users often call it 'the algorithm' or 'the bot'.
Recommender / feed	The system that decides which posts a user sees and in what order (e.g. a 'For You' feed). Reducing a post's place here is a common moderation action short of removal.
Bot	An automated account or program. Can be benign (a helper) or malicious (fake accounts used to harass or manipulate).
Decision-support / human-in-the-loop	A design where the AI advises but a person decides. The framework's core stance: the steward is always the final check — the workshop's 'the last word is yours'.

Sources (accessed 8 July 2026): Standard AI / trust-and-safety usage; framework working paper and workshop metaphors (LPA). Definitions written for a non-technical audience.

02 Platform & moderation terms

The vocabulary of how platforms govern content, cross-referenced to Corpus Material 1. Knowing these helps a steward read a platform's decision and set realistic expectations.

Community Standards / Guidelines	A platform's rulebook for what content is allowed (Meta's Community Standards, TikTok's Community Guidelines, the X Rules). Distinct from the Terms of Service.
Report / flag	To notify the platform that content may break its rules. Some categories accept reports from any user; some require the targeted person to report.
Takedown / removal	Deleting content for breaking the rules. Across platforms this is less common for coded or borderline hostility than reduced visibility.
Visibility reduction / downranking	Showing content to fewer people without deleting it — e.g. making it ineligible for the 'For You' feed. 'Not recommended' is not the same as 'removed'.
Shadowban	Informal term for suspected silent visibility reduction the user is not told about. Often inferred rather than confirmed.
Demonetisation	Removing a creator's ability to earn money from content. A common lever on video platforms and a driver of algospeak.
Suspension / ban	Temporarily or permanently disabling an account. Reserved for repeat or severe violations.
Deplatforming	Removing a person or group's presence across a service (or services), usually for severe or repeated harm.
Appeal	Asking a platform to review a decision. Meta-family cases can escalate to the independent Oversight Board.
Block / mute / restrict / hide	User-side controls: block cuts off an account; mute hides it from you; restrict (Instagram) quietly limits it; hiding/keyword filters suppress specific words or comments. Often the most reliable protection when removal is uncertain.
Community Notes	Crowd-sourced context added to posts (used on X and, in the US, Meta) instead of, or alongside, formal fact-checking.
Terms of Service (ToS)	The contract governing use of a platform — accounts, rights, liability. Broader than, and separate from, the content rules.

Sources (accessed 8 July 2026): *Corpus Material 1 (Platform Terms & Policies)* and the platforms' own policy pages; standard trust-and-safety usage.

03 Online-harm & harassment tactics

Names for the behaviours a community moderation tool is meant to help with. Naming a tactic makes it easier to spot, explain and act on.

Hate speech	Content attacking people on the basis of a protected characteristic (race, religion, ethnicity, sexual orientation, disability, and so on). Definitions vary by platform and by UK law (see Corpus Material 1).
Harassment	Targeting a person with abuse, threats or unwanted contact. Can be one sustained account or many people at once.
Dehumanisation	Portraying people as less than human (as animals, disease, filth, objects). A serious escalator that platforms treat as high-severity.
Dog-whistle	Coded language that seems innocent to outsiders but signals a hostile meaning to an in-group, preserving deniability. See Section 04.
Pile-on / dogpiling	Many users attacking one target at once, so that each message seems minor but the aggregate is overwhelming.
Brigading	Coordinated group action to swarm a target or community — mass replies, mass reporting, or flooding a thread.
Coordinated inauthentic behaviour	Organised use of fake or duplicate accounts to manufacture consensus, harass, or manipulate what appears popular.
Sealioning	Feigning polite, sincere debate while bombarding someone with persistent demands for evidence, to exhaust and derail them.
Concern trolling	Posing as a sympathetic insider in order to undermine a person or cause from within.
‘Just asking questions’ (JAQing)	Spreading a hostile claim through loaded questions rather than assertions, to keep deniability.
Whataboutism	Deflecting criticism by pointing to someone else’s alleged wrongdoing instead of answering.
Astroturfing	Disguising an organised campaign as spontaneous grassroots opinion.
Sockpuppet	A fake secondary account used to praise oneself, evade a ban, or fake agreement.
Doxing	Publishing someone’s private details (address, workplace, family) to intimidate or enable harassment.
Swatting	Making a false emergency report to send armed police to a target’s home. Rare but dangerous.
Gaslighting	Manipulating someone into doubting their own memory or perception of events.
DARVO	‘Deny, Attack, and Reverse Victim and Offender’ — an abuser recasting themselves as the real victim.
Trolling / grieving	Provoking or disrupting for reaction or amusement, ranging from mischief to targeted cruelty.
Catfishing	Using a fake identity to deceive someone, often to manipulate or exploit.

Sources (accessed 8 July 2026): Standard trust-and-safety and online-harms literature; framework typologies. Definitions paraphrased.

04 How hostility hides: coded language & evasion

Before the specific examples in Section 05, this section explains the *mechanics* — the general techniques by which meaning is hidden from outsiders and from automated filters. Understanding the mechanism matters more than memorising any single term, because the specifics change constantly.

Dog-whistle	Language (or a symbol, number or image) that reads as innocent to the general public but carries a specific hostile meaning for an in-group. Its whole point is deniability, which is why intent and context are decisive.
Algospeak	Coded expressions adopted to evade automated moderation — e.g. 'unalive' for kill/suicide or 'seggs' for sex. Crucially dual-use : marginalised users rely on it to survive over-moderation, while bad actors use it to slip harmful content past filters. Context decides which.
Leetspeak / character substitution	Swapping letters for numbers or symbols (e.g. mixing digits into a word) to dodge keyword filters. Older and more easily caught than invented words.
Homoglyphs	Using look-alike characters from other alphabets so a word reads normally to humans but not to filters.
Emoji / symbol substitution	Using emojis or symbols to stand in for words — sometimes playful, sometimes to encode a slur or a coded signal while evading detection.
Evaded slur	A slur written to escape filters (asterisks, spacing, misspellings, look-alikes). Their function is precisely to escape detection, including in the datasets platforms train on — which is why the community-curated evasion register exists.
Voldemorting	Deliberately not naming a person, brand or term (describing it instead) to avoid triggering search, tagging or detection.
Reclaimed / in-group use	A term that would be a slur from an outsider but is used affectionately or defiantly within the community it refers to. The classic over-moderation trap: automated tools remove it, misreading in-group speech as attack.

The two-sided lesson of coded language

The same techniques serve opposite ends. A queer creator writing 'le dollar bean', or a community member reclaiming a term, is protecting themselves from blunt filters; a hostile actor using a number code or a coded slur is hiding an attack. A moderation tool that treats all coded language as hostile will silence exactly the communities it is meant to protect. This is why every entry that follows is a prompt to look at context, speaker and pattern — not a verdict.

Sources (accessed 8 July 2026): *Algospeak research and reporting* (Steen, Yurechko & Klug 2023; Aleksic 2025; Washington Post; Wikipedia 'Algospeak'); *dog-whistle scholarship* (OUPblog; academic studies). Paraphrased for recognition.

05 Recognising coded hostility (for these communities)

A curated, non-exhaustive set of coded forms documented by established anti-hate monitors, prioritising the hostilities the partner communities report. These are **recognition aids**: a match means ‘look closer and consider context’, never ‘this person is hateful’. The live, current specifics belong in the evasion register and in the specialist databases named below.

Responsible-use rules for this section

(1) Context and intent decide meaning; many listed forms have entirely innocent everyday uses. (2) Never act on a single signal or accuse an individual from a match alone; look for pattern, target and intent. (3) These are for protecting the community, not for policing ordinary speech. (4) This list ages quickly — treat it as a teaching sample and keep the evasion register current.

Anti-migrant & anti-refugee

Coded form	What it can signal	Note for moderators
‘Illegals’ / ‘bogus asylum seekers’	Framing asylum seekers as inherently criminal or fraudulent, stripping their humanity.	Criticising immigration <i>policy</i> is legitimate; labelling <i>people</i> as illegitimate is the dehumanising move to watch for.
‘Invasion’ / ‘flooding’ / ‘swarms’	Water- and pest-metaphors that cast migrants as a natural disaster or infestation.	Metaphor choice is the signal; these frames recur just before spikes in real-world hostility.
‘Culture enricher’ (sarcastic)	A term of neo-Nazi origin used sarcastically to blame migrants — especially Muslims — for crime.	Sarcasm is the tell; used approvingly it signals a coded racist frame.
‘Great Replacement’ / ‘remigration’	A conspiracy claim that native populations are being deliberately replaced; ‘remigration’ euphemises mass expulsion.	Well-documented extremist frames; their appearance in ordinary debate is a strong flag.

Antisemitic

Coded form	What it can signal	Note for moderators
‘Globalists’ / ‘international financiers’	Long-standing coded stand-ins for a conspiracy about Jewish control of finance or politics.	Has innocent economic uses; the conspiracy framing (hidden puppet-masters) is what distinguishes the dog-whistle.
Triple parentheses (‘echoes’)	Placing a name inside ((())) to mark someone as Jewish, often as a target.	A recognised harassment marker; treat as a hostile flag.
‘Cultural Marxism’	A conspiracy trope with antisemitic roots blaming a hidden elite for social change.	Context-dependent, but a common bridge from mainstream grievance to extremist framing.
Holocaust denial framed as ‘jokes’ or questions	Minimising or denying the Holocaust through humour or ‘just asking’ formats.	The joke/question wrapper is the deniability device; treat the underlying claim as the harm.

Anti-Muslim

Coded form	What it can signal	Note for moderators
Racialised 'grooming gangs' framing	Weaponising a genuine safeguarding term to smear Muslim or South Asian men collectively.	Child protection is real and serious; the coded move is collective ethnic/religious blame. Handle with care and escalate.
'No-go zones' / 'creeping Sharia'	Claims that Muslim communities are lawless enclaves or covertly imposing law.	Recurring anti-Muslim frames; flag when used to portray a whole community as a threat.
Terror-conflation	Treating ordinary Muslims as presumptively linked to terrorism.	The generalisation from individuals to a whole faith is the tell.

Anti-Black

Coded form	What it can signal	Note for moderators
Coded criminality / dependency framing	Language (e.g. 'welfare' and 'thug' tropes) implying Black people are criminal or a drain, without saying so outright.	Historic, well-documented dog-whistle pattern; watch for the implied racial subject.
Decontextualised 'crime statistic' memes	Numbers stripped of context and posted to imply inherent Black criminality.	A recognised format; the missing context is deliberate. Treat as a hostile frame, not data.

Anti-East/Southeast-Asian (ESEA)

Coded form	What it can signal	Note for moderators
Disease / 'bat-eating' / COVID-blame tropes	Associating ESEA people with disease, dirtiness or blame for the pandemic.	The most-reported coded hostility for this community; often muted and coded post-pandemic (see Corpus Material 2).
State-conflation ('CCP' / 'spy')	Treating individuals of Chinese heritage as agents or proxies of the Chinese state.	Criticism of a government is legitimate; casting people as suspect by heritage is the harm.
'Go back to...' / perpetual-foreigner	Framing ESEA people as forever foreign regardless of how long settled.	A recurring exclusionary frame; flag the 'you don't belong' move.

Number & symbol codes (recognition only)

Coded form	What it can signal	Note for moderators
'88', '14', '1488'	Established white-supremacist number codes ('88' = 'HH'; '14' = a 14-word slogan).	Extremely context-dependent — also appear innocently (dates, sports numbers). Never judge on the number alone; weigh the whole account and context.
Co-opted memes (e.g. 'Clown World', certain frog/'honk' memes)	Some mainstream memes have been adopted as xenophobic/neo-Nazi in-jokes.	Co-option is partial; most uses are ordinary. Look for clustering with other signals.

Sources (accessed 8 July 2026): Recognition definitions synthesised from established anti-hate education: OUPblog on dog-whistles; University of Northampton extreme-right explainer; Political Research Associates glossary; academic dog-whistle research; and, for live tracking, the standard monitors stewards should consult — ADL 'Hate on Display', Hope not Hate, Community Security Trust (antisemitism) and Tell MAMA (anti-Muslim hate). No extremist sources are cited or linked. Graphic slurs and slogans are deliberately not reproduced.

06 Popular internet slang & shorthand

Everyday terms a moderator meets that are usually not hostile in themselves but shape tone and meaning. Included so the agent reads context correctly and does not mistake ordinary slang for harm.

Ratio	When a reply gets far more engagement than the post it answers — a sign the crowd disagrees. Can tip into pile-on.
Based / cringe	'Based' = admirably unafraid to hold a view; 'cringe' = embarrassing. 'Based' is sometimes used approvingly of edgy or bigoted takes — context matters.
Cope / seethe / mald	Taunts implying an opponent is upset and in denial. Common in pile-ons.
NPC	'Non-player character' — mocking someone as thoughtless or a follower; can be dehumanising in political use.
Touch grass	Telling someone to go offline / get perspective. Usually mild mockery.
IYKYK	'If you know, you know' — signals in-group meaning outsiders won't catch; sometimes shields coded content.
/s	Marks the preceding text as sarcasm. Its absence is a common source of misreading.
TW / CW	'Trigger warning' / 'content warning' — a heads-up before sensitive content.
OP	'Original poster' — the person who started the thread.
DM / slide into the DMs	'Direct message'; private messaging, sometimes a harassment channel.
QRT / dunk	'Quote-retweet' used to mock; 'dunking' is publicly ridiculing a post to one's own audience — can trigger pile-ons.
Milkshake duck	Someone embraced online then quickly exposed as problematic.

Sources (accessed 8 July 2026): General internet-culture usage; definitions paraphrased and kept brief.

07 Using the glossary in moderation

How the agent and stewards should apply this reference. The glossary informs judgement; it does not replace it.

As a router look-up

- When a steward asks what a term means, answer from the glossary in plain language and cite this entry.
- When the corpus does not contain a term, say so plainly rather than guessing, and offer to have a steward add it — new slang appears constantly.

As a prompt-level terminology behaviour

- Use the terms the community is comfortable with, and define any technical or unfamiliar term the first time it appears in a recommendation.
- When a coded form from Sections 04–05 may be present, name the possibility, explain the mechanism, and point to context — do not declare intent.

Boundaries

Do not treat a glossary match as proof of hostility; coded language is dual-use and context-dependent. Do not use the recognition sections to police ordinary speech or accuse individuals on a single signal. Do not decode or define terminology used to trade or access child sexual abuse material — such requests are refused, not answered. When meaning turns on community-specific or fast-changing knowledge the corpus lacks, surface the uncertainty and defer to a steward, consistent with the framework's uncertainty-surfacing and consent-to-escalate constraints.

Steward maintenance checklist

- Add new terms as they appear and retire ones that have fallen out of use; online language moves fast.
- Keep the recognition sections (04–05) synchronised with the community evasion register and with specialist monitors (ADL, Hope not Hate, CST, Tell MAMA).
- Flag any term whose meaning has shifted (e.g. a neutral word newly adopted as a code, or a code that has faded).
- Preserve the dual-use and context caveats on every coded entry when editing.
- Record a change history so the audit trail shows which community decision authorised each edit.

Provenance

AI, platform and harassment definitions reflect standard trust-and-safety usage and this corpus's companion materials. Algospeak entries draw on academic research and mainstream reporting. Coded-hostility recognition entries are synthesised from established anti-hate education and point stewards to specialist live databases; no extremist sources are cited, and graphic slurs and slogans are not reproduced. All definitions are paraphrased.

Disclaimer

This glossary is an educational reference to support community moderation, not a detection rule-set and not a verdict on any person. Coded language is dual-use; a term appearing here does not make its user hostile. Meaning depends on speaker, audience and context, which only a human steward can weigh. Online language changes continuously, so treat every entry as current only as of the access date.

End of Corpus Material 3 of 4 — Glossary of Tech & Popular Terms. Companions: (1) Platform Terms & Policies; (2) Political & Cultural Context. Next: (4) Persona categories.