

Platform Terms, Rules & Content-Moderation Policies

Facebook • Instagram • X (Twitter) • TikTok

Corpus reference document prepared for the **Layered Participatory Alignment (LPA)** community-in-the-loop moderation framework. It serves two audiences at once. For the **decision-support agent**, it supplies the platform-policy evidence needed to satisfy the framework's *retrieval-grounded explainability* and *platform-realism* constraints. For the **community stewards** at Race Council Cymru and the Chinese in Wales Association who operate and oversee the tool, it sets out, in plain terms, what each platform prohibits, how to report harms *from Wales*, what UK law adds when platform rules fall short, and how these map onto the harms their communities actually experience — including the coded anti-East/Southeast-Asian hostility that automated systems routinely miss.

Document	Platform Terms & Policies — Corpus Material 1 of 4
Maps to	Layer 1 retrieval corpus → Legal & regulatory reference index (High priority, 11 files)
Serves	The LPA decision-support agent and community stewards at Race Council Cymru and the Chinese in Wales Association
Jurisdiction	UK / Wales — platform rules read alongside the Online Safety Act 2023, Equality Act 2010 and Welsh hate-crime reporting routes
Feeds constraints	Platform-realism; retrieval-grounded explainability; consent-to-escalate
Coverage	Facebook / Meta, Instagram, X (Twitter), TikTok
Content date	Policies as published/accessed 8 July 2026
Status	Living document — re-verify against source URLs each review cycle

Currency warning — read before use

Platform moderation policies change frequently and were in unusual flux during 2025–2026. Meta renamed and substantially loosened its hate-speech rules in January 2025; X pared down its hateful-conduct policy after 2023; TikTok updated its Community Guidelines in the second half of 2025. Treat every statement here as accurate **as accessed on 8 July 2026** and confirm against the cited source URLs before relying on it operationally. Community stewards may edit this entry as policies shift.

Copyright note. All policy content below is paraphrased in the compilers' own words and attributed to its source. No platform policy text is reproduced verbatim. URLs are provided so stewards can consult the authoritative originals.

This is one entry in the community-curated retrieval corpus that grounds the LPA moderation agent. Within Step 4 of the framework, the agent must cite specific corpus entries when it explains a recommendation and must *refuse rather than guess* when retrieval returns nothing relevant. This document supplies the platform-policy evidence for that grounding: it lets the agent tell a community steward not only *whether* a piece of content appears to breach a platform rule, but *which* rule, *how* that platform tends to act on it, and *what* a realistic outcome of reporting would be.

Three framework ideas shape what is included and emphasised:

- **Platform-realism.** The system prompt requires recommendations calibrated to how platforms actually behave, acknowledging that takedowns are not always achievable. Each section therefore ends with an explicit ‘what is realistically achievable’ note rather than assuming reported content is removed.
- **Reclaimed and in-group language.** The framework’s central concern is over-moderation of reclaimed or in-group expression alongside under-moderation of coded hostility. Every platform’s treatment of context and reclaimed slurs is captured, because these carve-outs are exactly where automated tools fail the communities LPA serves.
- **Coded hostility / the evasion register.** The corpus maintains a community-curated evasion-pattern register for dog-whistles and evaded slurs that public datasets under-represent. Where a platform explicitly recognises coded hate (TikTok most clearly), that is flagged so the register can cross-reference platform language.

Consent-to-escalate. The agent never auto-routes to police or platform reporting without explicit human action. The reporting pathways below are reference material a human steward can act on — not triggers for automated reporting.

Who this document serves

The partner communities — Race Council Cymru (representing minority ethnic communities across Wales) and the Chinese in Wales Association — are not passive subjects of moderation but the governors of the tool. This document is written so a steward with no technical or legal background can use it directly: to understand which platform rule a piece of content engages, to judge how likely a platform is to act, to choose a Wales-based reporting route, and to see where UK law provides leverage that a platform’s own terms do not. Two dedicated sections at the end — *UK & Wales legal and reporting context* and *Mapping the communities’ harms to platform coverage* — exist for that purpose, with a specific focus on the coded anti-East/Southeast-Asian hostility the Chinese community reports.

Reading each platform entry. Sections follow a fixed structure so retrieval returns comparable chunks: *Governing documents* → *Protected characteristics* → *Prohibited hateful content* → *Harassment & bullying* → *Context & reclaimed language* → *Coded / evaded hostility* → *How to report* → *Enforcement & appeals* → *What is realistically achievable* → *Sources*.

High-level comparison of the four platforms on the dimensions most relevant to moderation support for minority ethnic communities. Cells are compressed summaries; see each platform section for detail and sources.

Dimension	Facebook / Meta	Instagram	X (Twitter)	TikTok
Core rulebook	Meta Community Standards (shared across apps)	Same Meta Community Standards + IG Community Guidelines	The X Rules + topic policies	TikTok Community Guidelines
Hate policy name	Hateful Conduct (renamed from Hate Speech, Jan 2025)	Hateful Conduct (via Meta)	Hateful Conduct policy	Hate Speech & Hateful Behavior
Reclaimed / in-group use	Allowed where speaker intent is clearly self-referential/empowering	Same as Meta	Consensual reclaimed slurs not treated as abuse when intent clear	Assessed via context & intent (captions, hashtags)
Coded-hate recognition	'Proxies' for protected traits; some cases need extra context	Same as Meta	Slurs/tropes graded by severity & repetition	Explicit: code words, symbols, audio trends, memes
Bystander reports	Accepted for many categories; some require target self-report	Same as Meta	Accepted (first-person & bystander); some need target input	Accepted; reporter identity kept confidential
Middle enforcement option	Reduced distribution / warning screens	Reduced distribution; comment controls	Limited visibility; labels; public-interest notice	Ineligible for For You feed (reduced reach)
Proactive enforcement	Scaled back in 2025 for less-severe violations; more report-led	Same as Meta	Often requires the target to report before action	Tech + human moderators; removal, restriction or ban
Appeals	In-app appeal + Oversight Board for eligible cases	In-app appeal + Oversight Board	In-app appeal of suspensions/actions	In-app appeal of removals/bans

Reading note for the agent. 'Not recommended' or 'reduced reach' is not the same as 'removed'. Across all four platforms, the most common outcome for borderline or coded content is reduced visibility rather than deletion, and on several platforms action depends on the target reporting. Recommendations should say so plainly.

Governing docs	Facebook Terms of Service; Meta Privacy Policy; Meta Community Standards (the substantive content rules); topic policies in the Transparency Center. UK users have a locale-specific Terms page.
Scope	The Community Standards apply to everyone worldwide and to all content, including AI-generated content. They also govern Instagram and Threads.
Structure	Each standard opens with a 'Policy Rationale', then lists content that is not allowed, and content that needs extra context to enforce, is allowed behind a warning screen, or is limited to adults.

Protected characteristics. Meta's Hateful Conduct standard protects direct attacks on people (not concepts or institutions) based on race, ethnicity, national origin, disability, religious affiliation, caste, sexual orientation, sex, gender identity and serious disease. Age counts when paired with another protected characteristic. Refugees, migrants, immigrants and asylum seekers get protection from the most severe (Tier 1) attacks, although criticism of immigration policy is permitted. Some non-protected traits (e.g. occupation) are covered when referenced together with a protected characteristic, and locally recognised proxy terms may be treated as standing in for a protected characteristic.

Prohibited hateful content. Meta says it removes dehumanising speech, allegations of serious immorality or criminality, and slurs, along with harmful stereotypes — dehumanising comparisons historically used to attack or exclude a group (its examples include Blackface, Holocaust denial, and tropes about Jewish control of institutions). It also removes serious insults, expressions of contempt or disgust, cursing, and calls for exclusion or segregation when aimed at people on the basis of a protected characteristic. Violations are split into two severity tiers.

Harassment & bullying. A separate Bullying and Harassment standard sets tiered protections. Everyone is protected from a baseline set of attacks (for example calls for death or serious disease, and certain sexualised degradation); minors, private adults and limited-scope public figures receive additional protection, including from dehumanising comparisons to animals, disease or filth and from attacks on physical appearance. For some categories the target (or an authorised representative) must report the content before Meta will act. These protections extend across pages, groups, events and messages.

Context & reclaimed language. Meta states it allows slurs or someone else's speech that are shared to condemn or report on them, and it allows self-referential or empowering use of language that would otherwise violate the rules where the speaker's intention is clear. A satire exception can also apply. This is the policy hook that in-group reclaimed expression relies on — and the point at which automated systems most often over-remove.

Coded / evaded hostility. Meta's main lever here is treating locally recognised proxy terms as protected-characteristic references, plus a 'needs additional context to enforce' tier for attacks on concepts or institutions tied to a protected group that are likely to contribute to imminent physical harm, intimidation or discrimination. Signals it weighs include potential to incite imminent violence, periods of heightened tension such as elections or conflict, and a recent history of violence against the targeted group.

How to report	Use the in-product report control on the post, comment, profile, group or message. EU users who believe content breaches a local hate-speech law can additionally file a legal (local-law) removal request.
Enforcement actions	Content removal; warning/sensitivity screens; reduced distribution (demotion); feature and account restrictions; account disabling for repeat or severe violations.
Appeals	In-product appeal of most decisions; eligible cases can be escalated to the independent Oversight Board, whose published decisions have reshaped several Meta policies.

What is realistically achievable (platform-realism)

In January 2025 Meta announced it would concentrate proactive automated enforcement on illegal and high-severity violations and rely more on user reports for less-severe ones — a category that in practice includes much implicit hate and bullying. It also replaced third-party fact-checking with Community Notes in the US. Practical implication for recommendations: coded or 'lawful-but-awful' hostility is **less likely to be removed proactively** than it once was; a user report is now often necessary; and reclaimed in-group language, though protected on paper, still risks automated over-removal. Blocking, muting and reduced-distribution outcomes are realistic even where deletion is not.

Sources (accessed 8 July 2026): Meta Community Standards index & Policies overview (transparency.meta.com/policies); Hateful Conduct (transparency.meta.com/policies/community-standards/hateful-conduct); Bullying and Harassment (transparency.meta.com/policies/community-standards/bullying-harassment); Community Standards Enforcement & Integrity Reports 2025 (transparency.meta.com/reports); Facebook Terms (facebook.com/legal/terms, facebook.com/terms/?locale=en_GB); Meta Privacy Policy (facebook.com/privacy/policy). Meta's Jan 2025 policy change is widely documented by third parties (e.g. GLAAD, HRC, CCDH).

Governing docs	Instagram Terms of Use; Meta Community Standards (the same substantive rules as Facebook and Threads); Instagram's own Community Guidelines, which restate those rules in shorter form; Meta Privacy Policy.
Relationship to Meta	Instagram does not maintain a separate hate-speech or harassment rulebook. Its prohibited-content definitions — protected characteristics, hateful conduct tiers, bullying protections, reclaimed-language carve-outs — are the Meta Community Standards summarised in the Facebook / Meta section above.

What is Instagram-specific. Instagram layers additional user-side anti-harassment tooling on top of the shared standards. Relevant controls include Hidden Words (filtering offensive comments and message requests, including custom word/phrase lists that can hold a community's known slurs and coded terms), Limits (temporarily restricting interactions during a spike of abuse), Restrict (quietly limiting a specific account), and granular comment controls. For a community steward, these tools are often a faster, more reliable mitigation than waiting on a content-removal decision.

Reporting. Report content, comments, accounts or messages in-app through the report control. As on Facebook, some bullying-and-harassment categories require the target (or an authorised representative) to report before Instagram will act, while hate-speech reports may be filed by bystanders.

What is realistically achievable (platform-realism)

Because Instagram runs on the same enforcement posture as Meta, the same 2025 realism applies: proactive removal is concentrated on illegal and high-severity content, and implicit or coded hostility often needs a user report. The practical upside on Instagram is the strength of its **user-side controls** (Hidden Words, Limits, Restrict): recommendations can reliably point a steward to filtering and interaction limits even when content removal is uncertain. Note also that Instagram limits the algorithmic reach of political content, which can affect how community-facing advocacy posts perform.

Sources (accessed 8 July 2026): Meta Community Standards (transparency.meta.com/policies/community-standards) as applied to Instagram; Instagram Community Guidelines / Help Center (help.instagram.com/581066165581870); Instagram Terms of Use; Meta Privacy Policy (facebook.com/privacy/policy).

Governing docs	X Terms of Service; X Privacy Policy; The X Rules (safety, privacy, authenticity) plus individual topic policies — Hateful Conduct, Abusive Behavior, Violent Speech, Violent & Hateful Entities, and Abusive Profile.
Structure	The X Rules give one-line statements of each prohibition, each linking to a fuller policy page. Enforcement philosophy and the range of enforcement options are documented separately in the Help Center.

Protected characteristics. Under its Hateful Conduct policy, X says users may not attack others on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability or serious disease.

Prohibited hateful content. The policy prohibits several categories aimed at people because of a protected characteristic: hateful references to violence or violent events where that group was the primary target, shared with intent to harass; repeated slurs and tropes intended to degrade or reinforce negative stereotypes; dehumanisation of a group; and hateful imagery or symbols (its examples include the Nazi swastika). A separate Abusive Profile rule bars using a username, display name or bio for targeted harassment or to express hate toward a protected category.

Harassment & bullying. The Abusive Behavior policy prohibits targeting people with abuse or harassment, or encouraging others to do so. It covers malicious, unreciprocated targeting (for example repeated replies or mass mentions/tags meant to humiliate), accounts dedicated to harassing someone, coordinated pile-ons, and denial of a verified mass-casualty event when shared with abusive context.

Context & reclaimed language. X states that posts can look hateful in isolation but not within a larger conversation, and that when members of a protected group use reclaimed terms consensually the context is not abusive. Its teams say they sometimes need to hear directly from the targeted person before acting, and that criticism of institutions, practices and ideas is protected expression rather than a violation.

Coded / evaded hostility. X grades slur-and-trope enforcement by severity and repetition rather than naming coded terms: severe or repetitive slur use in a harassing context can require post removal, while moderate or isolated use may instead have its visibility limited. There is no published coded-term register, so the community-curated evasion register is especially important when reasoning about X.

How to report	Anyone can report from a post, profile or Direct Message (first-person or bystander); X may ask the reporter to select additional posts for context. For violent threats, the reporter can email themselves a copy of the report to share with law enforcement. A confirmation may take up to 24 hours.
Enforcement actions	Post-level: limiting visibility/discoverability, labelling, a public-interest notice, or requiring removal (often with a read-only period). Account-level: read-only mode, ownership verification, or suspension. Penalty depends on severity, whether a person was targeted, and prior record.
Immediate suspension	Egregious violations — e.g. violent threats, non-consensual intimate media, child sexual exploitation — can result in immediate and permanent suspension.
Appeals	Users who believe an account was actioned or suspended in error can submit an appeal.

What is realistically achievable (platform-realism)

X's hateful-conduct rules were significantly pared down after the 2023 ownership change (the ban on incitement to violence against protected groups was moved into the violent-speech policy), and enforcement leans heavily on **the targeted person reporting** and on **context**. Practical implications: for coded or moderate hostility, the realistic outcome is often reduced visibility rather than removal; action frequently requires the target (not just a bystander) to report; and the public-interest exception can keep otherwise-violating posts online behind a notice. Recommendations should treat removal as uncertain and foreground blocking, muting and report-with-context as reliable steward actions.

Sources (accessed 8 July 2026): *The X Rules* (help.x.com/en/rules-and-policies/x-rules); *Hateful Conduct policy* ([.../hateful-conduct-policy](https://help.x.com/en/rules-and-policies/hateful-conduct-policy)); *Abusive Behavior* ([.../abusive-behavior](https://help.x.com/en/rules-and-policies/abusive-behavior)); *Abusive Profile* ([.../abusive-profile](https://help.x.com/en/rules-and-policies/abusive-profile)); *Violent Entities* ([.../violent-entities](https://help.x.com/en/rules-and-policies/violent-entities)); *Enforcement options* ([.../enforcement-options](https://help.x.com/en/rules-and-policies/enforcement-options)) and *Enforcement philosophy* ([.../enforcement-philosophy](https://help.x.com/en/rules-and-policies/enforcement-philosophy)); *Report abusive behavior* (help.x.com/en/safety-and-security/report-abusive-behavior); *Rules Enforcement report* (transparency.x.com/en/reports/rules-enforcement). Policy history per independent analyses of X's 2023 changes.

Governing docs	TikTok Terms of Service (region-specific — separate US and EEA versions, the latter with a plain-language summary); TikTok Community Guidelines (the substantive content rules, updated in the second half of 2025); Privacy Policy; and Safety Center explainers.
Structure	Community Guidelines are grouped into themes. The ‘Safety & Civility’ theme carries the rules most relevant here: Hate Speech & Hateful Behavior, Violent & Hateful Organizations and Individuals, and Harassment & Bullying.

Protected attributes & prohibited hateful content. TikTok says it does not allow hate speech, hateful behavior, or promotion of hateful ideologies, including content that attacks people based on protected attributes such as race, religion, gender or sexual orientation. It states it will not recommend content containing negative stereotypes about a protected group even where the content is not removed, and — to account for regional and cultural sensitivities — may reduce visibility or remove content that disparages religion in regions where the risk of harm is greater. Hateful ideologies (belief systems that exclude, oppress or discriminate based on protected attributes) are barred, including associated names, symbols, slogans, gestures, music or imagery, and claims of inherent group superiority or inferiority.

Harassment & bullying. TikTok prohibits harassment and bullying, including in retaliation: degrading remarks about appearance, doxing, sexual harassment, wishing serious harm or death on someone, threats, depicting someone being bullied, and inciting harassment or coordinated abuse (for example urging others to leave abusive comments or to file false reports). It allows more critical commentary about public figures than private ones, but still removes the most harmful content (doxing, sexual harassment, threats, hate). Lower-severity degrading content may simply be made ineligible for the For You feed.

Context & reclaimed language. TikTok assesses harassment using context, intent and whether conduct is coordinated, and uses indicators such as captions and hashtags to gauge intent — the mechanism through which consensual or educational use is distinguished from an attack.

Coded / evaded hostility — strongest explicit recognition of the four

TikTok’s hate guidance explicitly names the tactics the LPA evasion register targets: it states that hate can be conveyed through images, text, audio, cartoons, memes, objects, gestures and symbols — ‘some of which won’t always be obvious’ — and specifically lists coded messages used by hate groups (code words, symbols, audio trends), jokes and memes that normalise discrimination, and hate-founded conspiracy theories. This makes TikTok the most useful platform to cross-reference when building and justifying evasion-register entries.

How to report	Report content or accounts in-app (e.g. Hate and harassment → Harassment and bullying paths, with options for yourself, someone you know, or a public figure). TikTok says it does not disclose the reporter’s identity to the reported party and keeps the reporter informed of the report’s status.
Safety tools	Controls over who can comment, Duet, Stitch or message you; private accounts; comment filters, keyword filters and mute — usable to hold a community’s known coded terms.
Enforcement actions	A combination of automated technology and human moderators. Outcomes range across removal, ineligibility for the For You feed (reduced reach), and account bans. Violent and hateful organisations or individuals are reviewed — including off-platform behavior — and may be banned.
Appeals	Users can appeal content removals and account actions in-app.

What is realistically achievable (platform-realism)

TikTok relies heavily on a middle option — making content **ineligible for the For You feed** (cutting reach) rather than removing it — so 'not recommended' is a common and realistic outcome for borderline or coded content, distinct from deletion. Its policy language on coded hate is strong, but detection of the subtlest coded material is imperfect, which is precisely why the community-curated evasion register adds value.

Recommendations can reliably point stewards to reporting plus TikTok's keyword-filter and interaction controls as dependable mitigations.

Sources (accessed 8 July 2026): *TikTok Community Guidelines — Overview and Safety & Civility* (tiktok.com/community-guidelines/en/safety-civility; tiktok.com/safety/en/policies-and-engagement/overview); *Countering hate speech & behavior and Bullying prevention* ([tiktok.com/safety/en/community-support/...](https://tiktok.com/safety/en/community-support/)); *Policies & engagement / content moderation* (tiktok.com/safety/en/policies-and-engagement); *Terms of Service* (tiktok.com/legal/page/us/terms-of-service/en and [.../global/summary-of-terms-and-conditions-eea/en](https://tiktok.com/legal/page/global/summary-of-terms-and-conditions-eea/en)).

A platform's own rules are only half the picture. The partner communities are in Wales, so the platforms operate here under UK law, and stewards have Welsh reporting routes that exist regardless of whether a platform acts. This section gives stewards (and the agent) that backstop: what UK law requires of platforms, where platform rules end and the law begins, and exactly where to report hate from Wales.

The Online Safety Act 2023 and Ofcom

The Online Safety Act 2023 (OSA) received Royal Assent in October 2023, and its illegal-content duties came into force on 17 March 2025, regulated by Ofcom. In-scope user-to-user and search services — which include all four platforms in this document — must assess the risk of illegal content and take proportionate steps to remove it quickly once aware and to reduce the risk of 'priority' illegal content appearing. The priority-offence list covers, among other things, terrorism, child sexual abuse material, and offences of stirring up racial and religious hatred. The largest (Category 1) services also have to enforce their own terms of service consistently and transparently — so a platform that says it removes hateful content can, in principle, be held to that promise.

Ofcom has opened investigations into dozens of services and can levy fines of up to £18 million or 10% of global turnover, with its stated 2026 priorities including tackling illegal hate and protecting women and girls online. For a steward, this means a platform's inaction on clearly illegal hate is now a regulatory matter, not only a policy preference.

Why the OSA does not close the gap LPA addresses

The OSA bites mainly on content that is **illegal** under UK law. Much of what these communities experience — coded slurs, dog-whistles, long polite-sounding arguments that a group is 'incompatible with British society' — is lawful-but-harmful and falls **outside** the illegal-content duties, landing back on each platform's own terms and transparency obligations. That is precisely the space community-curated moderation is built to cover: the law handles the illegal tail, the platforms under-detect the coded middle, and the community corpus and evasion register fill the gap.

Equality Act 2010 and hate-crime law

The UK's baseline for protected characteristics is the Equality Act 2010: age, disability, gender reassignment, marriage and civil partnership, pregnancy and maternity, race, religion or belief, sex, and sexual orientation. Platform 'protected-characteristic' lists overlap with this but are not identical (for example X and Meta add caste), so a steward should not assume a platform rule and UK law draw the line in the same place. For hate crime specifically, the police monitor five strands — race, religion, sexual orientation, disability and transgender identity — and hostility based on these can aggravate an offence and lead to a stronger sentence. Online abuse can be a hate crime, but not every hurtful post is criminal; the law targets threats, harassment and stirring-up-hatred offences rather than offensiveness alone.

In Wales, this sits alongside the Well-being of Future Generations (Wales) Act 2015, whose goals include a more equal Wales and a Wales of cohesive communities, and the Welsh Government's anti-racist ambition and its 'Hate hurts Wales' campaign — context stewards can draw on when pressing for action or support.

Where to report hate from Wales (steward reference)

These routes work whether or not a platform removes the content, and several allow a report on behalf of someone else or anonymously — useful when a community member does not want to report themselves.

Police	999 in an emergency; 101 (or online) for non-emergencies. Online abuse and social-media hate can be reported.
True Vision / True Vision Cymru	The police-run hate-crime reporting site (report-it.org.uk). Online form that reaches your local force; supports reporting hate seen online or on social media, and reporting on behalf of someone else.
Wales Hate Support Centre	Run by Victim Support; free, confidential, in English and Welsh; 0300 3031 982. An independent alternative to the police that can make a third-party report to the police on your behalf, offers an anonymous option, and supports victims and witnesses whether or not they report to the police.
Specialist services	Tell MAMA (anti-Muslim hate); Community Security Trust / CST (antisemitism); Stop Hate UK; Crimestoppers (anonymous, 0800 555 111).
Community networks	besea.n (Britain's East and Southeast Asian Network) and allied ESEA networks for the Chinese / ESEA community; Race Council Cymru as a Wales-wide community anchor.

Scale in Wales. Victim Support’s Welsh hate-crime service recorded about 2,230 referrals in 2024–25, up roughly 13% year on year, with over half race-based — a reminder that race is the dominant hate strand in Wales and that demand is rising, even though much online hostility is never reported at all.

Sources (accessed 8 July 2026): *Online Safety Act 2023* (legislation.gov.uk/ukpga/2023/50) and GOV.UK OSA explainer; Ofcom *illegal-content duties and enforcement updates* (ofcom.org.uk/online-safety). *Equality Act 2010*; Home Office *hate crime statistics, England & Wales, year ending March 2025* (gov.uk). *Well-being of Future Generations (Wales) Act 2015*; ‘Hate hurts Wales’ (gov.wales/hate-hurts-wales). True Vision / True Vision Cymru (report-it.org.uk/wales); Wales Hate Support Centre — Victim Support (reporthate.victimsupport.org.uk); Victim Support Wales referral figures 2024–25.

This section connects the platform rules above to the harms the partner communities actually report, so a steward or the agent can see, for each harm pattern, which platform rule applies, how platforms tend to act, and how to respond from Wales. Harm patterns draw on the framework's Step 1–2 findings and on documented experiences of minority ethnic communities in the UK, including the Chinese / East and Southeast Asian (ESEA) community.

Harm pattern the community reports	Platform rule that applies	How platforms tend to act	Respond / report
Explicit racial slurs & dehumanising attacks (race / ethnicity)	Prohibited under every platform's hate policy; slurs and dehumanising comparisons named explicitly	Most likely category to be removed; explicit slurs are the clearest case	Report in-app; True Vision / Wales Hate Support Centre; may be a hate crime
Anti-Muslim & anti-religious hostility, often implicit / discursive (religion / belief)	Prohibited as hate on protected grounds; but usually needs context to enforce	Under-detected by automated tools — the framework's analysis found implicit religious hate frequently scored low-severity	Report in-app + Tell MAMA; use keyword filters; steward judgement
Anti-migrant / refugee & 'incompatible with British society' narratives (national origin)	Attacks on people prohibited, but criticism of immigration policy is allowed — a fine line	Long, polite-sounding arguments often scored low-severity and left up; Meta may act faster during heightened tension (e.g. around the 2024 UK riots)	Escalate to a human; report in-app; document pattern
Coded anti-East/Southeast-Asian hostility — disease / 'bat-eating' / COVID-blame tropes, conflation with the Chinese state, model-minority erasure	Sits in the gap: often lawful-but-harmful; TikTok's coded-hate language is the closest fit	Frequently missed by automated detection; reduced reach is more likely than removal	Lean on the community evasion register; report + filter; besea.n / ESEA support
Reclaimed / in-group language read as harmful out of context	All four protect consensual reclaimed use where intent is clear	Automated tools over-remove exactly this — an over-moderation, not under-moderation, failure	Invoke the carve-out; appeal wrongful removals
Pile-ons / coordinated harassment (mass replies, tagging, brigading)	Prohibited under harassment / abusive-behavior policies; coordinated abuse named on Meta, X, TikTok	Often requires the targeted person to report; individual posts may each seem minor	Block / mute / Limits; report the pattern; may be a hate crime

Focus: the Chinese in Wales / ESEA community

The Chinese and wider East and Southeast Asian community reports a distinctive band of hostility that public datasets and automated classifiers under-represent: disease- and COVID-blame framing and 'bat-eating' tropes, the pairing of ESEA faces with negative stories, conflation of a diverse community with the Chinese state, and a structural invisibility that hides the problem — ESEA people are often folded into an aggregated 'Asian: other' category, so the harm is neither counted nor named. Under-reporting is high: during the pandemic, hate crimes against Chinese people in London roughly tripled early on, and a 2020 poll found a large majority of British Chinese respondents had received racial slurs. Although those figures are pandemic-era, the coded *pattern* persists. Because so much of it is coded and lawful-but-harmful, it is exactly the content the community-curated corpus and evasion register — not the platforms' automated systems — are best placed to catch. Practical implication for the agent: weight community-curated evidence heavily for this community, expect reduced-reach rather than removal, and surface Wales reporting routes and besea.n / ESEA support.

Sources (accessed 8 July 2026): Framework Step 1–2 findings (LPA working paper); documented ESEA / British Chinese experiences and pandemic-era hate-crime rises (e.g. openDemocracy; TIME; University of Nottingham pilot study; peer-reviewed research on ESEA young people's experiences). Community organisations: besea.n and allied ESEA networks. Wales reporting routes as cited in the previous section.

These observations synthesise the four platform entries against the framework's design commitments. They are written to be retrieved alongside the platform sections when the agent reasons about a case.

Reclaimed and in-group language (over-moderation risk)

All four platforms carve out consensual, self-referential or reclaimed use of otherwise-violating language when the speaker's intent is clear — Meta via self-referential/empowering use, X via consensual reclaimed slurs, TikTok via context-and-intent indicators, and Instagram via Meta's standards. Yet the framework's own analysis and the platforms' enforcement records show automated systems still over-remove exactly this speech. The corpus should therefore treat a policy carve-out as necessary but not sufficient: the agent should surface the carve-out, flag the over-moderation risk, and leave the final judgement to the community steward.

Coded hostility, dog-whistles and the evasion register

Platform coverage of coded hostility is uneven. TikTok is the most explicit, naming code words, symbols, audio trends and memes. Meta relies on locally recognised proxy terms plus a needs-more-context tier. X grades slur-and-trope enforcement by severity and repetition without a published coded-term list. The gap between platform recognition and real-world coded hostility is the space the community-curated evasion register fills; where a platform's own language matches a register entry (most often TikTok's), the agent can cite it to strengthen an explanation.

Platform-realism: what reporting actually achieves

- **Removal is not the default outcome.** The most common result for borderline or coded content is reduced visibility (Meta demotion, X visibility limits, TikTok For-You ineligibility), not deletion.
- **Reporting is often a precondition.** Meta (post-2025) leans on user reports for less-severe violations; X frequently needs the targeted person specifically to report; TikTok and Meta accept bystander reports for hate but require target reports for some harassment categories.
- **Exceptions keep content up.** X's public-interest / newsworthiness exception and Meta's satire and needs-context tiers can leave otherwise-borderline content online.
- **User-side controls are the reliable lever.** Blocking, muting, comment/keyword filters, Instagram Limits and TikTok interaction controls give stewards dependable mitigations even when removal is uncertain.

The system prompt's platform-realism constraint should draw on these points so recommendations never imply a takedown is guaranteed.

Consent-to-escalate and legal-removal channels

The reporting pathways in this document are reference material for a human steward, consistent with the framework's rule that the agent never routes to platform reporting or police without explicit human action. Section 06 sets out the concrete Wales routes — police (999 / 101), True Vision Cymru, and the Wales Hate Support Centre, plus specialist services such as Tell MAMA — that a steward can choose from. Two further legal channels sit alongside the platforms' own rules: EU/EEA users can file local-law removal requests with Meta, and UK duties under the Online Safety Act 2023 operate independently of any single platform's community standards. The agent should treat platform rules, UK law, and Welsh reporting routes as three distinct grounds for action, surface all that apply, and leave the choice to the steward.

Steward maintenance checklist

- Re-verify each platform section against the cited source URLs at the start of every review cycle; record the new access date.
- Watch for renamed or restructured policies (Meta renamed Hate Speech to Hateful Conduct in 2025; X moved incitement provisions between policies in 2023).
- Log material changes to enforcement posture (proactive vs report-led) because these drive the platform-realism notes more than the prohibition wording does.
- Cross-check TikTok's coded-hate language and any new proxy-term guidance against the evasion register.
- Re-verify the UK / Wales context: Online Safety Act enforcement stage, Ofcom priorities, and that the Wales reporting routes (numbers, services) are still current, since these change too.
- Keep the community-harm mapping current as the partner communities report new coded patterns; treat the Chinese / ESEA focus box as editable by those stewards.
- Keep a change history so the audit trail shows which community decision authorised each edit.

Provenance

Compiled from the platforms' official Terms, Rules, Community Standards/Guidelines and Transparency Center / Safety Center pages, supplemented by reputable third-party documentation of 2023–2026 policy changes where those changes post-date or contextualise the official pages. The UK / Wales context draws on primary legal sources (the Online Safety Act 2023, Equality Act 2010, Well-being of Future Generations (Wales) Act 2015), Ofcom and GOV.UK / GOV.WALES material, and Welsh reporting services (True Vision Cymru, the Wales Hate Support Centre). The community-harm mapping draws on the LPA working paper's Step 1–2 findings and documented ESEA / British Chinese experiences. All content is paraphrased; source URLs are listed at the foot of each section for verification.

Disclaimer

This document is a reference summary for community moderation support, not legal advice and not an official statement of any platform's policy. Platform rules and their enforcement change frequently; the authoritative source is always the platform's own current policy pages. Where a decision has legal consequences, consult the primary legal instruments and a qualified adviser.

End of Corpus Material 1 of 4 — Platform Terms & Policies. Next: (2) Political and cultural context; (3) Glossary of tech and popular terms; (4) Persona categories.