

Persona Categories & Personas

Viewpoints the agent uses to read content and to shape replies

A set of structured viewpoints for community moderation in Wales. Each persona is used **twice**: first as an **analysis lens** — a viewpoint the agent reasons through to interpret a piece of content (how would this land for an isolated elder, a targeted member, or a hostile actor?) — and second as a **response lens** — a viewpoint the agent tailors its reply to (plain and reassuring for one audience, concise and platform-aware for another). The personas span both sides of the moderation encounter: the **community members and stewards who use the tool**, and the **harm sources whose content it must handle**.

Document	Persona Categories & Personas — Corpus Material 4 of 4
Used twice	(1) Analyse incoming content from multiple viewpoints; (2) generate responses tailored to a viewpoint
Covers	Community members & audiences; community-member profiles by origin; operators & stewards; harm sources; easily-misread viewpoints; political-leaning viewpoints
Serves	The LPA decision-support agent and community stewards at Race Council Cymru and the Chinese in Wales Association
Companions	(1) Platform Terms; (2) Political & Cultural Context; (3) Glossary
Status	Draft archetypes for community validation — not real people, not fixed descriptions

Two boundaries built into these personas

1. Harm-source personas are analysis-only. The agent uses them to *understand intent and recognise hostility* when reading content — never as a voice to write in, and never to help such an actor. The response lens applies only to community members, audiences and stewards. **2. Personas are archetypes, not people.** They are composite viewpoints drawn from documented community context to aid perspective-taking; they are internally diverse and must never be used to stereotype a real person or infer someone’s identity, politics or intent from a single cue.

These personas operationalise the framework’s central insight: the same words carry opposite meanings depending on who uses them, to whom, and within which community. Rather than one ‘view from nowhere’, the agent reasons through several situated viewpoints. Each persona therefore does two distinct jobs.

Use 1 — Analysis: reading content from multiple viewpoints

When content arrives (incoming) or is being checked before posting (pre-publication), the agent considers it through the relevant personas: how a targeted member receives it, how an isolated elder might miss or be distressed by it, what a coded-hostility actor may intend by it, whether it is in-group reclaimed speech that only looks hostile. The output is a richer interpretation and an explicit note of where viewpoints diverge — handed to the steward, never decided unilaterally.

Use 2 — Response: replying according to a viewpoint

When the agent drafts a recommendation, a suggested reply, or a pre-publication check, it tailors register, language, length and emphasis to the audience persona — plain, gentle and step-by-step for an elder; concise and platform-aware for a digitally-fluent youth; reasoning-rich and options-based for a steward; dual-audience for a community communicator. **This response lens is used only for community-side personas.**

Why harm-source personas never drive responses

The agent models a harm source only to *analyse* — to gauge intent, recognise tactics and assess target and severity. It does not write in a harm-source voice, produce content that serves that actor’s aim, or help refine hostile messaging. This keeps the personas consistent with the framework’s refusal to generate malicious content and with child-safety limits, while still letting the agent understand what it is looking at.

01 Persona quick reference

The twenty-four personas across six categories, with what each contributes to analysis and to response. Full cards follow in Sections 02–07.

Persona	Category	Analysis use	Response use
Community elder	Members	May miss coded hostility; distress; scam risk	Plain, gentle, multilingual, step-by-step
Second/third-gen youth	Members	Reads in-group / reclaimed use; platform-fluent	Concise, direct, platform-aware
Recently-arrived migrant	Members	Alert to diaspora politics; unfamiliar with UK routes	Neutral, reassuring, orient to UK routes
Targeted member	Members	The 'how the target receives it' test	Trauma-aware; safety options first
Mei-Ling (Chinese)	Origin profile	Lives on WeChat/RedNote; conflict-avoidant; visa stress	Mandarin framing; private, low-pressure
Emmanuel (Black)	Origin profile	Church network; structural-racism attuned; Welsh-admin gap	Community-oriented; practical; non-tokenising
Harry (South Asian)	Origin profile	Trusted broker; bridges audiences; elder-care strain	Respect standing; nuanced; bilingual
Joy (Filipino)	Origin profile	Time-poor NHS shifts; anti-migrant-worker hostility	Concise, low-effort, wellbeing-aware
Tariq (Gulf)	Origin profile	Attuned to Islamophobia; student; digital-dependent	Religiously aware; neutral; ISoc support
Steward / manager	Operators	Weights legitimacy, risk, precedent	Cited reasoning + options, not directives
Frontline caseworker	Operators	Time-pressured triage of volume	Prioritised, actionable, batch summaries
Community communicator	Operators	Tests wording against two audiences	Dual-audience read; suggest, don't rewrite
Coded-hostility actor	Harm – analysis only	Detect dog-whistles / evaded slurs by pattern	— (not embodied)
Coordinated harasser	Harm – analysis only	Spot pile-ons / brigading as a pattern	— (not embodied)
'Polite' bad-faith actor	Harm – analysis only	Name sealioning / exclusionary premise	— (not embodied)
Opportunistic pile-on	Harm – analysis only	Contextualise surges in tension	— (not embodied)
Disinformation amplifier	Harm – analysis only	Flag false / defamatory, high-reach harm	— (not embodied)
Reclaimed-language user	Easily misread	Guard against over-moderation of in-group speech	Protect the speech; ask steward if unsure
Legitimate critic	Easily misread	Policy/state criticism ≠ attack on people	Protect debate; don't treat as hate
Progressive / justice	Political	Low harm threshold; names harm explicitly	Take harm seriously; avoid over-flagging
Liberal / free-expression	Political	Higher threshold; wary of over-removal	Proportionality; explain reasoning

Persona	Category	Analysis use	Response use
Conservative / traditionalist	Political	Legitimate traditional views ≠ hate	Neutral, respectful; don't mislabel
Immigration-sceptic	Political	Policy critique vs attack-on-people line	Engage neutrally; flag only if it crosses
Politically disengaged	Political	May miss political coding / dog-whistles	Plain, non-partisan, explain neutrally

'Harm – analysis only' personas have no response use by design: the agent reasons about them to interpret content, but never writes in their voice.

Viewpoints of the people the tool ultimately serves and protects. Used for analysis (how content lands for them) and for response (how to address them well).

Community elder (first generation)		Members · analysis + response
Who they are	A long-settled first-generation member, often Cantonese- or Hakka-speaking, with limited English and digital confidence and a raised risk of isolation (the Chinese in Wales Association's core group). Relies on the community organisation for advice.	
As an analysis lens	May not recognise algospeak or coded hostility; may under-report; sarcasm and nuance can be missed across a language barrier; heightened exposure to scams and impersonation; content others shrug off can cause real distress.	
For responses	Plain language and short sentences; offer Welsh- or Chinese-language framing; define every technical term (link the glossary); be gentle and reassuring; give concrete, step-by-step help and a human contact; never assume digital fluency.	
Cross-refs	Context §04 (isolation, language); Glossary (define terms); Platform Terms (reporting routes; user controls).	

Second/third-generation youth (British/Welsh-born)		Members · analysis + response
Who they are	Bilingual and digitally fluent, moving fluently across platforms; navigating a dual Welsh and minority-ethnic identity; encounters reclaimed language and coded hostility routinely.	
As an analysis lens	Recognises in-group and reclaimed use and algospeak; can usually tell banter from attack; may use reclaimed terms themselves (an over-moderation risk); attuned to TikTok/Instagram-native coded hostility.	
For responses	Can take concise, direct, platform-aware answers; respect their reading of in-group language and protect it from over-flagging; don't over-explain the basics.	
Cross-refs	Glossary §04–05 (coded language); Context §05 (reclaimed language).	

Recently-arrived migrant (Hong Kong BN(O) exemplar)		Members · analysis + response
Who they are	A recent arrival — e.g. on the Hong Kong BN(O) route — often Cantonese-speaking, in politically sensitive circumstances, navigating unfamiliar UK systems under a 'no recourse to public funds' condition. (Generalises to other new arrivals.)	
As an analysis lens	Highly alert to political content and diaspora-politics fault lines; may read state-conflation or surveillance themes as threatening; unfamiliar with UK platforms and reporting routes; vulnerable to targeted disinformation.	
For responses	Strict neutrality on diaspora politics; never infer a political stance from heritage; orient to UK routes; offer multilingual framing; be careful and honest about confidentiality rather than giving absolute guarantees.	
Cross-refs	Context §04–05 (BN(O); diaspora politics; state-conflation); Platform Terms (reporting).	

Targeted member (recipient of hostility)		Members · analysis + response
Who they are	A community member directly on the receiving end of hostile content — the viewpoint whose safety and wellbeing are most immediately at stake.	
As an analysis lens	The essential 'how does the target receive this?' test: coded or implicit hostility a classifier scores low can be experienced as unmistakably hostile by the person aimed at.	
For responses	Trauma-aware and validating without amplifying distress; lead with safety and options (block/mute/report/support); never minimise; surface support resources; don't make them relive detail unnecessarily.	
Cross-refs	Platform Terms (user controls; realism); Context (community harms); wellbeing resources.	

Sources (accessed 8 July 2026): Community context and needs from Corpus Material 2 (Political & Cultural Context) and the partner organisations' described services. Personas are illustrative archetypes, not real individuals.

Category 2 — Community-member profiles (origin & lived experience)

Five richer, lived-experience profiles that complement the lighter interaction archetypes of Category 1. Each is a composite drawn from documented community experience across Wales, spanning different origins, locations and structural barriers. Like every persona here they are used twice — as an analysis lens and, because they are community members, as a response lens — and they remain archetypes for perspective-taking, never descriptions of real people.

Persona	Primary cultural hub	Main geographic zone	Core systemic barrier
Mei-Ling (Chinese)	Academic / digital (WeChat)	Cardiff (urban south)	Visa anxiety & local civic disconnect
Emmanuel (Black)	Faith community (church)	Newport (M4 corridor)	Promotion ceilings & school language barriers
Harinder (South Asian)	Religious / business (Gurdwara)	Swansea (urban west)	Business strain & culturally-competent elder care
Joy (Filipino)	Professional diaspora (NHS)	Bangor (rural north)	Rural isolation & occupational burnout
Tariq (Gulf)	University / religious (ISoc)	Aberystwyth (coastal mid)	Alcohol-centric spaces & dietary scarcity

Mei-Ling Zhou · 34 · Chinese

Origin profile · analysis + response

Snapshot	Biostatistics postdoctoral researcher at Cardiff University, living in Cardiff Bay. Moved from Shenzhen seven years ago for her PhD; well integrated academically but socially bound to international networks.
Lifestyle & routines	Long, irregular lab and home-desk hours; walks the Bay barrage; batch-cooks Cantonese and Sichuan food from Asian supermarkets on Cowbridge Road East; socialises mainly with Chinese academics and expatriates; streams mainland dramas; occasional London trips.
Behaviour & traits	Professional and polite in English but most herself in Mandarin; avoids voicing grievances publicly (conflict-avoidance). Lives in a digital parallel world — WeChat for roughly 90% of personal communication, news and organising, and RedNote (Xiaohongshu) for lifestyle and reviews rather than Instagram or TikTok.
Pain points	Visa and job insecurity (funding / sponsorship anxiety); systemic isolation inside a diverse but insular university bubble; the emotional toll of being an only child far from ageing parents in China.
Motivations & goals	A permanent lectureship or industry R&D role that secures long-term residency; a way to bridge her cultural identity with Cardiff civic life without feeling a perpetual outsider.
As an analysis lens	Much of her online life is on WeChat and RedNote — Chinese-language platforms outside the four covered in Corpus Material 1 and outside Latin-script classifiers; she may under-report and avoid public complaint; highly alert to visa / immigration content and to China / diaspora-politics framing.
For responses	Offer Mandarin framing; keep it private and low-pressure (respect conflict-avoidance); recognise her platforms differ from the mainstream four; be sensitive to visa stress; never assume she uses Western platforms.
Cross-refs	Context §04–05 (Chinese community; diaspora politics); Platform Terms (note: WeChat / RedNote sit outside this corpus's four platforms); Glossary.

Emmanuel Achebe · 42 · Black (British-Nigerian)		Origin profile · analysis + response
Snapshot	Logistics manager at a distribution centre near the M4, in Newport. Relocated his family from south-east London four years ago for affordable housing and a safer, slower pace to raise three young children.	
Lifestyle & routines	Early shifts at the logistics hub; evenings and weekends given over to family logistics (football, music lessons, school); anchored by a vibrant multicultural Pentecostal church in Newport; Saturday gatherings with Jollof rice and pounded yam and Premier League football.	
Behaviour & traits	Community-minded ('it takes a village'), actively building a neighbourhood network; a cultural preservationist, fiercely proud of his Igbo heritage, who keeps the language alive at home while supporting his children's local integration.	
Pain points	The Welsh-language gap — his children thrive in Welsh-medium schools, but he and his wife cannot help with homework or read school admin; subtle workplace bias, with senior promotions bypassing him toward long-standing local networks; and the loss of London's dense Black-led cultural spaces, networks and services (e.g. barbershops for Afro-textured hair).	
Motivations & goals	The best educational foundation for his children without losing their cultural identity; breaking into South Wales professional networks to advocate for diversity in management.	
As an analysis lens	Attuned to structural and workplace-racism framing; his church network is a key information and support channel; likely to encounter anti-Black coded hostility (Glossary §05).	
For responses	Acknowledge structural concerns without tokenising; community- and church-oriented framing lands well; be practical about the Welsh-language admin barrier; keep it concise given his time pressures.	
Cross-refs	Context §02 (Newport; communities); Glossary §05 (anti-Black coded hostility); Platform Terms.	

Harinder 'Harry' Singh · 58 · South Asian (Sikh)		Origin profile · analysis + response
Snapshot	Independent community-pharmacy owner in Swansea. His parents migrated from Punjab in the late 1980s; he grew up in the city, trained in Cardiff, and runs the family business; a prominent Sikh community elder and Gurdwara trustee.	
Lifestyle & routines	Works six days a week serving a diverse, working-class demographic; primary carer for his 82-year-old mother, who lives with him, his wife and his youngest adult son; Sunday mornings for Seva and worship at the Gurdwara; follows Indian politics and cricket.	
Behaviour & traits	A trusted advisor — authoritative yet empathetic — whom locals of all backgrounds treat as a community counsellor; an effortless code-switcher across English, Punjabi and local Welsh phrases.	
Pain points	Small-business strain (overheads, changing NHS pharmacy contracts, supply-chain issues, little council support); culturally insufficient elder care that understands his mother's strict-vegetarian diet, language and religious needs; and a generational disconnect with UK-born children less drawn to the business or traditional structures.	
Motivations & goals	Keeping the pharmacy viable to pass on or sell for retirement; better representation and structural support for ethnic-minority elders within NHS Wales care.	
As an analysis lens	A trusted community broker whose reading carries weight; attuned to anti-Sikh and anti-South-Asian hostility and to religious-dress abuse; naturally bridges white-Welsh and South-Asian audiences.	
For responses	Respect his standing as an elder and community leader; he can handle nuanced, policy-level content; support his bridging role; use bilingual framing where it helps.	
Cross-refs	Context §02, §05 (religion as a distinct axis); Glossary; Platform Terms.	

Maria 'Joy' Santos · 29 · Filipino		Origin profile · analysis + response
Snapshot	Senior staff nurse (Band 6) at Ysbyty Gwynedd (NHS Wales), in Bangor. Recruited from Manila four years ago in an international NHS drive; shares a terraced house with three other Filipino nurses.	
Lifestyle & routines	A grueling roster of 12.5-hour day and night shifts; days off spent sleeping, video-calling family in the Philippines across a 7–8 hour time difference, and cooking Filipino comfort food; her social circle is the Filipino healthcare diaspora in North Wales, with Sunday church and group trips to Chester or Liverpool.	
Behaviour & traits	Resilient and duty-driven, with an ironclad work ethic rooted in responsibility to family; a financial minimalist who tracks every penny around the monthly remittance she sends for her siblings' tuition.	
Pain points	Extreme rural isolation (the shift from Manila to Bangor, dark wet winters, poor transport without a car); professional burnout amid NHS staffing shortages, feeling relied upon for labour but rarely fast-tracked; and a double-identity barrier, navigating English charts while many elderly patients speak Welsh.	
Motivations & goals	To sponsor a younger sibling into UK healthcare and share the family's financial weight; to gain specialised training toward higher-paying urban or specialist roles.	
As an analysis lens	Time-poor from long shifts, with little bandwidth for lengthy exchanges; alert to anti-migrant-worker and NHS-related hostility ('taking our jobs'); connects digitally through video calls and diaspora networks.	
For responses	Concise and low-effort; wellbeing-aware (burnout, isolation, seasonal strain); practical and quick, respecting her limited time.	
Cross-refs	Context §01–02 (Nation of Sanctuary; migrant workers); Glossary §03 (harassment); Platform Terms.	

Tariq Al-Mansoor · 23 · Gulf (Omani)		Origin profile · analysis + response
Snapshot	Final-year Renewable Energy Engineering undergraduate at Aberystwyth University, on a full Omani Ministry of Higher Education scholarship, in university accommodation.	
Lifestyle & routines	Lectures, engineering labs and study at the National Library of Wales; basic self-catering; gym and casual football; an active member of the University Islamic Society (ISoc); weekends in common rooms, gaming online with friends in Oman, and walking the promenade.	
Behaviour & traits	Culturally reserved — polite, observant, quiet in group settings until comfortable, then warm; abstains from alcohol, which shapes his social boundaries; tech-dependent for daily video calls, banking, food delivery and forums.	
Pain points	An alcohol-centric student culture that makes it hard to bond with domestic students without drinking; practical and religious scarcity (consistent Halal options need bulk orders from larger cities; acute homesickness at Ramadan and Eid); and seasonal affective strain from harsh, low-light coastal winters during exam season.	
Motivations & goals	A First or Upper-Second to meet his scholarship requirements and make his family proud; a short-term UK green-energy graduate placement before returning to Oman.	
As an analysis lens	Attuned to anti-Muslim / Islamophobic content and mockery of religious practice; navigating life as an international student; digital-dependent; may be targeted around religious holidays or dress.	
For responses	Religiously aware and neutral; orient to student and ISoc support; sensitive to Islamophobia; concise and practical.	
Cross-refs	Context §05 (religion; Islamophobia); Glossary §05 (anti-Muslim coded hostility); Platform Terms.	

Sources (accessed 8 July 2026): Detailed profiles incorporate community-partner persona research supplied for this project, tightened and aligned to the corpus format, with lived-experience context cross-referenced to Corpus Material 2. Composite archetypes only — not real individuals, and containing no personal data.

The people who actually run and oversee the tool. Used for analysis (what they need to judge) and for response (how the agent should serve them). The steward is always the final decision-maker.

Steward / manager (the final check)		Operators · analysis + response
Who they are	A senior staff member who operates and oversees the tool and makes the final decision, accountable to the community and to partners.	
As an analysis lens	Weighs legitimacy ('will colleagues accept this?'), organisational risk and precedent; wants the reasoning, not just a verdict; embodies the framework's 'final check'.	
For responses	Give inspectable reasoning with cited corpus entries; present options and trade-offs rather than directives; support defensibility to colleagues; surface uncertainty; never auto-act (consent-to-escalate).	
Cross-refs	All corpus materials; framework governance (retrieval-grounded explainability; consent-to-escalate).	

Frontline caseworker / engagement officer		Operators · analysis + response
Who they are	Handles day-to-day incoming hostility and high-volume triage and supports members directly; usually time-pressured.	
As an analysis lens	Needs fast triage grouped by type and severity; separates the urgent (threats, doxing) from noise; lives the high-volume pile-on context surfaced in the framework's Step 1.	
For responses	Concise, prioritised, actionable summaries; batch triage where volume is high; clear next steps and escalation paths; realistic about what reporting will and won't achieve.	
Cross-refs	Router (batch/triage mode); Platform Terms (realism); Glossary (harm tactics).	

Community communicator (dual-audience)		Operators · analysis + response
Who they are	Drafts outgoing public statements that must speak to government partners and community members at once — the pre-publication branch made human.	
As an analysis lens	Tests wording against two audiences; alert to how a statement could be misread by one side or weaponised by hostile actors.	
For responses	Provide a dual-audience read (government and community), name where the two diverge, flag weaponisation risk, and consider bilingual reception; suggest rather than rewrite wholesale — final wording rests with them.	
Cross-refs	Context §05 (the bridge tension); framework pre-publication branch.	

Sources (accessed 8 July 2026): Task contexts and the decision-support / final-check principle from the LPA working paper (Steps 1 and 4–5).

Analysis-only — read before using this category

These are recognition lenses for understanding intent and detecting hostility when the agent *reads* content. The agent must never write in these voices, produce content that advances their aims, or help refine hostile messaging. Recognition cues below are deliberately kept at the level of pattern-spotting, not how-to. Every assessment is a prompt to look at target, pattern and context and to escalate to a steward — never a verdict on an individual from a single cue.

Coded-hostility actor / dog-whistler

Harm source · ANALYSIS ONLY

What they do	Attack while preserving deniability, using dog-whistles, algospeak, evaded slurs or coded symbols so the hostility is legible to an in-group but not to outsiders or blunt filters.
Recognition cues	Coded forms from Glossary §05; several signals clustering together; a consistent target; 'plausible deniability' framing. A single coded term alone is not enough.
How the agent responds	Surface the possible coded reading and the mechanism, cite the evasion register / glossary, assess target and pattern, recommend action per platform realism, and escalate — without declaring intent.
Cross-refs	Glossary §04–05; evasion register; Platform Terms (realism).

Coordinated harasser

Harm source · ANALYSIS ONLY

What they do	Drive or take part in coordinated pile-ons, brigading, mass-reporting or dedicated harassment accounts.
Recognition cues	A volume spike; synchronised timing or near-identical phrasing; coordination across accounts; mass mentions or tags aimed at one target.
How the agent responds	Triage the campaign as a single pattern rather than post-by-post; advise user-side controls (block/mute/Limits) plus reporting the pattern; note possible hate-crime dimension; escalate.
Cross-refs	Glossary §03 (brigading, pile-on); Platform Terms (harassment; user controls); UK/Wales routes.

'Polite' bad-faith actor

Harm source · ANALYSIS ONLY

What they do	Exclude or exhaust while staying within civility rules — sealioning, 'just asking', concern trolling, or long polite-sounding arguments that a group does not belong.
Recognition cues	Superficial politeness over an exclusionary premise; relentless demands framed as debate; the framework's 'long polite-sounding argument' scenario.
How the agent responds	Name the tactic and the underlying exclusionary move; separate it from legitimate debate (see the Legitimate critic persona); support a member's choice to disengage; be realistic that platforms often will not act.
Cross-refs	Glossary §03 (sealioning, JAQing); Context §06 (anti-migrant narratives).

Opportunistic pile-on participant		Harm source · ANALYSIS ONLY
What they do	Not organised, but joins in during heightened tension (after a triggering event), amplifying hostility in the moment.	
Recognition cues	Clustering around news or events; low-effort echoing of a dominant line; emotionally driven rather than premeditated.	
How the agent responds	Contextualise the surge as a heightened-tension period; advise volume-management (mute/limit); triage by severity; distinguish a transient spike from sustained targeting.	
Cross-refs	Context §06 (current climate; post-2024 dynamics); Platform Terms (visibility, controls).	

Disinformation amplifier		Harm source · ANALYSIS ONLY
What they do	Spread false or defamatory content — sometimes at high reach — that drives real-world harm, as in the Welsh Refugee Council case in Corpus Material 2.	
Recognition cues	Unverified or fabricated claims; defamatory framing; high-reach amplification; targeting of an organisation or named individual.	
How the agent responds	Flag as possible disinformation or defamation; advise documenting and reporting, and — for organisations — escalating to human comms/legal support; note the online-to-offline risk; avoid amplifying through engagement.	
Cross-refs	Context §06 (Welsh Refugee Council case); Platform Terms (reporting, legal-removal); UK routes.	

Sources (accessed 8 July 2026): Harm-actor tactics synthesised from Corpus Material 1 (platform harms), Corpus Material 3 (harassment tactics & coded language) and Corpus Material 2 (current climate). Recognition-level only; no operational detail.

Why this category exists

These are the viewpoints automated tools most often get **wrong**, by reading in-group or legitimate speech as hostile. The framework's whole purpose includes preventing that over-moderation, so the agent's default here is to **protect the speech** and, where genuinely unsure, ask the steward rather than flag.

In-group reclaimed-language user

Easily misread · analysis + response

Who they are	A community member using reclaimed or in-group language that can look hostile out of context but is affectionate, defiant or ordinary within the group.
As an analysis lens	Speaker is in-group and intent is not hostile — the classic false positive. Platform policies explicitly protect consensual reclaimed use where intent is clear.
For responses	Default to protecting the speech; invoke the reclaimed-use carve-outs; do not flag as harm; if genuinely uncertain about speaker or intent, ask the steward rather than over-remove.
Cross-refs	Glossary §04 (reclaimed use); Platform Terms (carve-outs); Context §05.

Legitimate critic

Easily misread · analysis + response

Who they are	Someone criticising a policy, institution or government (including a foreign state) — not attacking people by who they are.
As an analysis lens	Criticism of ideas, policy or a state is not an attack on people or heritage — the state-conflation guard from the other side. This is protected expression across platforms and in law.
For responses	Don't treat policy or state criticism as hate; distinguish the target (an idea vs a person); protect legitimate debate; flag only if it crosses into attacking people on a protected characteristic.
Cross-refs	Context §05 (state vs people); Platform Terms (criticism of institutions allowed).

Sources (accessed 8 July 2026): Over-moderation concern and reclaimed-language / criticism carve-outs from the LPA working paper, Corpus Material 1 (platform policy) and Corpus Material 2 (sensitivities).

A spectrum of political viewpoints held across the wider public and within the communities themselves. Their purpose is to help the agent anticipate how the same content is read very differently across the spectrum (analysis) and to engage each audience fairly (response). They are audience lenses, not harm sources.

Neutrality — read before using this category

Political leaning is not hostility. These personas represent a range of legitimate viewpoints; none is treated as inherently suspect, more trustworthy, or a proxy for harm. Harm is judged on behaviour and tactics (Category 4), never on political identity. The agent stays politically neutral, never infers a real person's politics from a cue, and never uses these lenses to label or profile anyone. They exist partly to guard against a real risk in the other direction: mistaking legitimate political disagreement — especially views that criticise immigration policy or hold socially conservative positions — for hate.

Progressive / social-justice-oriented

Political lens · analysis + response

Who they are	Prioritises anti-racism, inclusion and structural change; often uses activist vocabulary; wants platforms to take harm seriously and to weigh impact and power, not just intent.
As an analysis lens	Tends to a lower threshold for harm and may read borderline or coded content as clearly hostile; attentive to structural context and to how content lands for those with least power.
For responses	They expect the tool to name and take harm seriously; the agent should do so while still distinguishing genuine harm from over-flagging and holding the same evidentiary bar for everyone.
Cross-refs	Category 4 (harm judged on behaviour); Category 5 (easily-misread guards); Glossary \$05.

Liberal / free-expression centrist

Political lens · analysis + response

Who they are	Values both inclusion and free expression; wary of silencing and over-reach; wants proportionate, clearly-explained responses.
As an analysis lens	Applies a higher threshold for harm and is alert to over-moderation and to legitimate speech being removed.
For responses	Emphasise proportionality and transparent reasoning; show why an action is or is not warranted rather than asserting it.
Cross-refs	Category 5 (easily-misread); Platform Terms (proportionate enforcement; appeals).

Conservative / traditionalist

Political lens · analysis + response

Who they are	Values tradition, faith, family and social order; may hold socially conservative views on religion, gender or national identity; concerned about being unfairly labelled.
As an analysis lens	Expressing traditional or religious views is legitimate expression, not hate; guard against mislabelling moral or religious disagreement as hostility.
For responses	Neutral and respectful; do not treat conservative or faith-based views as inherently suspect; apply the same rules as to anyone.
Cross-refs	Category 5 (Legitimate critic); Context \$05 (religion as a distinct axis).

Immigration-sceptic / populist		Political lens · analysis + response
Who they are	Critical of immigration levels and of multiculturalism as policy; may support parties campaigning on those issues.	
As an analysis lens	The point where the legitimate-critic vs coded-hostility line is most active: criticising immigration policy is legitimate, while dehumanising people or attacking them by a protected characteristic is not. Distinguish the two carefully rather than treating the whole viewpoint as suspect.	
For responses	Engage the policy argument neutrally; do not conflate the leaning with hostility or treat the person as a harm source; flag only where content crosses into attacking people. Compare against the Legitimate critic and the 'polite' bad-faith actor to tell them apart.	
Cross-refs	Category 4 ('polite' bad-faith actor); Category 5 (Legitimate critic); Context §06 (contested immigration politics).	

Politically disengaged / apolitical		Political lens · analysis + response
Who they are	Not driven by politics; uses platforms socially; lower political-media literacy.	
As an analysis lens	May share or repeat coded content without recognising it, and may miss dog-whistles entirely; their intent is often simply unaware rather than hostile.	
For responses	Plain and non-partisan; explain any political coding neutrally, without assuming motive.	
Cross-refs	Glossary §04–05 (coded language; recognition); Category 5 (easily-misread).	

Sources (accessed 8 July 2026): Political-leaning archetypes written to represent a fair spectrum of legitimate viewpoints, consistent with the framework's neutrality constraint. They are interpretive aids, not political judgements, and must never be used to label real individuals.

08 Using personas in moderation

How the agent puts the two uses into practice. Personas inform judgement; the steward decides.

Use 1 — multi-viewpoint analysis

- For incoming content, read it through the relevant viewpoints — at least the targeted member, plus any harm-source lens that might explain intent and any easily-misread lens that might exonerate it.
- Name where viewpoints diverge (e.g. ‘a classifier scores this low, but for the targeted member it reads as a clear threat’) instead of collapsing them into one verdict.
- Treat a harm-source match as a reason to look at target, pattern and context — never as proof of intent.
- Use the political-leaning lenses to anticipate how the same content reads across the spectrum — not to judge the author’s politics, and never to treat a viewpoint as hostile on its own.
- Hand the interpretation, with its uncertainties, to the steward.

Use 2 — viewpoint-tailored responses

- Identify the audience persona and match register, language, length and emphasis to it (plain and gentle for an elder; concise and platform-aware for a youth; reasoning-rich for a steward; dual-audience for a communicator).
- Define unfamiliar terms and offer bilingual framing where the audience needs it; draw on the origin profiles to match register and channel (for example, recognising that some members live on Chinese-language platforms, or that a time-poor shift worker needs brevity).
- For a suggested reply to hostile content, write in the community member’s protective voice — never in the hostile actor’s.

Boundaries

Never write in, or optimise content for, a harm-source persona. Never infer a real person’s identity, politics or intent from a persona match or a single cue — personas are archetypes, not detectors. Apply the political-leaning lenses to anticipate reception and engage fairly, never to treat any viewpoint as inherently hostile; harm is judged on behaviour, not politics. Default to protecting easily-misread speech. When a judgement turns on community-specific knowledge the corpus lacks, surface the uncertainty and escalate to a steward, consistent with the framework’s uncertainty-surfacing and consent-to-escalate constraints.

Steward maintenance checklist

- Validate each persona against lived community knowledge; add, merge, split or remove them — stewards are the authority, and the Chinese in Wales Association should shape the Chinese-community viewpoints.
- Keep archetypes from hardening into stereotypes; revise wording that could be read as a fixed description of real people — this applies especially to the detailed origin profiles.
- Keep the political-leaning lenses balanced and even-handed; no viewpoint should be written as a caricature or treated as a proxy for harm.
- Add new harm-source patterns as they emerge, keeping them analysis-only and recognition-level.
- Keep response guidance aligned with how community members actually prefer to be addressed.
- Record a change history so the audit trail shows which community decision authorised each edit.

Provenance

Personas are composite archetypes synthesised from this corpus — the community and organisational context of Corpus Material 2, the platform and harm landscape of Corpus Material 1, and the terminology and tactics of Corpus Material 3 — together with the task contexts and design commitments of the LPA working paper. The detailed origin profiles (Category 2) incorporate persona research supplied by the community partners, tightened and aligned to the corpus format; the political-leaning lenses (Category 6) are written to represent a fair spectrum of legitimate viewpoints. They are not drawn from real individuals and contain no personal data.

Disclaimer

These personas are perspective-taking aids to support community moderation, not descriptions of real people and not detection rules. Communities are internally diverse; no archetype captures any individual. A persona match is never evidence about a real person's identity or intent. The authoritative account of any community viewpoint is that community itself, through its stewards, who may revise these personas at any time.

End of Corpus Material 4 of 4 — Persona Categories & Personas. Completes the set with (1) Platform Terms & Policies, (2) Political & Cultural Context, and (3) Glossary of Tech & Popular Terms.